

Deep-Learning Capstone — Full Package

Every group builds the same system — an **Agentic-RAG agent over scanned books** — **differentiated** across **three axes (domain × data speciality × primary NFR)** so no two projects converge.

Edu Bungla + Math Calibration

Read in this order

1. **02-How-To-Submit** — how you hand work in (a GitHub repo, tagged at each milestone). Read this first.
2. **03-Project-Specification** — the whole system: pipeline, the three axes, the regulated-repo design, tiers, anti-copying.
3. **05-Codebase-Guide** — Part 1 documents every file; Part 2 is the ordered **57-step build** checklist. You build against this.
4. **forms/** — A2–A4, the remaining milestone forms you fill in (Word). A1 is already submitted.
5. **the repo itself** — the code you build in (`doc-agent-starter.zip`). Every file you touch is a named stub — **you edit, never create.**
6. **06-Group-Assignment-Workbook** — domains, data specialities, NFRs, data sources, and the build buckets (one workbook).

How the milestones work

You get ONE repo (from the starter template) and submit it **four times** by pushing a tag — `a2-submit`, `a3-submit`, `a4-submit`. Each milestone is the previous one plus your new work; your commit history is your submission. See **02-How-To-Submit** for the exact commands.

The single grading folder

Your repo's `grading_kit/` is the one folder that makes your project reproducible and gradable: `manifest.yaml` declares your three axes (domain, data speciality, primary NFR + its target) and points to your corpus and held-out pages. It is also your own evaluation harness.

Formats: read-only references are PDF; the forms stay editable Word; the Group-Assignment Workbook stays an interactive spreadsheet; the code stub is a repository.